

# Homework 3: Ensemble Methods and Recurrent Neural Networks

Large-Scale Data Analysis 2020

Shahriyar Mahdi Robbani (XQR418)

May 27, 2020

## 1 AdaBoost

Let  $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\} \in (\mathbb{R}^d \times \{-1, 1\})^n$  be the training data,  $t = 1, \dots, T$  denote the boosting rounds. At round  $t$ ,  $w_i^{(t)}$  is the weight of training pattern  $i$ ,  $h^{(t)}$  is the base classifier from round  $t$ ,  $\alpha^{(t)}$  is the importance of this classifier,  $\varepsilon^{(t)} = \sum_{j=1}^n w_j \mathbb{I}(h^{(t)}(\mathbf{x}_j) \neq y_j)$ , and the aggregated classifier is  $f^{(t)} = \sum_{i=1}^t \alpha^{(i)} h^{(i)}$ .

### 1.1 Step 1

The goal is to prove for each  $i$  at any round  $t$

$$w_i^{(t+1)} = w_i^{(t)} \frac{\exp(-\alpha^{(t)} h^{(t)}(\mathbf{x}_i) y_i)}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}}. \quad (1)$$

We start from the update of the weights as defined in the lecture:

$$w_i^{(t+1)} = \begin{cases} w_i^{(t)} / (2\varepsilon^{(t)}) & \text{if } h^{(t)}(\mathbf{x}_j) \neq y_j \\ w_i^{(t)} / (2(1 - \varepsilon^{(t)})) & \text{otherwise} \end{cases} \quad (2)$$

And importance  $\alpha^{(t)}$  of each weight defined in the lecture as:

$$\alpha^{(t)} = \frac{1}{2} \ln \left( \frac{1 - \varepsilon^{(t)}}{\varepsilon^{(t)}} \right) \quad (3)$$

The prediction  $h^{(t)}(\mathbf{x}_j)$  and class  $y_j$  can only take the values 1 and -1. Therefore when  $h^{(t)}(\mathbf{x}_j) = y_j$ , we can rewrite equation (1) as:

$$w_i^{(t+1)} = w_i^{(t)} \frac{\exp(-\alpha^{(t)})}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}} \quad (4)$$

Substituting equation (3) into equation (4) we get:

$$\begin{aligned} w_i^{(t+1)} &= w_i^{(t)} \frac{\exp\left(-\frac{1}{2} \ln\left(\frac{1 - \varepsilon^{(t)}}{\varepsilon^{(t)}}\right)\right)}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}} \\ &= w_i^{(t)} \frac{\left(\frac{1 - \varepsilon^{(t)}}{\varepsilon^{(t)}}\right)^{-\frac{1}{2}}}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}} \\ &= w_i^{(t)} \frac{\sqrt{1 - \varepsilon^{(t)}}}{\sqrt{\varepsilon^{(t)}}} \frac{1}{2\sqrt{\varepsilon^{(t)}}\sqrt{1 - \varepsilon^{(t)}}} \\ &= w_i^{(t)} / (2\varepsilon^{(t)}) \end{aligned} \quad (5)$$

Similarly, when  $h^{(t)}(\mathbf{x}_j) \neq y_j$ , we can rewrite equation (1) as:

$$w_i^{(t+1)} = w_i^{(t)} \frac{\exp(\alpha^{(t)})}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}} \quad (6)$$

Substituting equation (3) into equation (6) we get:

$$\begin{aligned} w_i^{(t+1)} &= w_i^{(t)} \frac{\exp\left(\frac{1}{2} \ln\left(\frac{1-\varepsilon^{(t)}}{\varepsilon^{(t)}}\right)\right)}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}} \\ &= w_i^{(t)} \frac{\left(\frac{1-\varepsilon^{(t)}}{\varepsilon^{(t)}}\right)^{\frac{1}{2}}}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}} \\ &= w_i^{(t)} \frac{\sqrt{\varepsilon^{(t)}}}{\sqrt{1 - \varepsilon^{(t)}}} \frac{1}{2\sqrt{\varepsilon^{(t)}}\sqrt{1 - \varepsilon^{(t)}}} \\ &= w_i^{(t)} / (2(1 - \varepsilon^{(t)})) \end{aligned} \quad (7)$$

We can see equations (5) and (7) are the two cases of equation (2), we have proved that equation (1) equals equation (2).

## 1.2 Step 2

The goal is to prove for each  $i$  at any round  $t$

$$w_i^{(t)} \frac{\exp(-\alpha^{(t)} h^{(t)}(\mathbf{x}_i) y_i)}{2\sqrt{\varepsilon^{(t)}(1 - \varepsilon^{(t)})}} = \frac{\exp(-f^{(t)}(\mathbf{x}_i) y_i)}{n \prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1 - \varepsilon^{(\tau)})}}. \quad (8)$$

Since each weight depends on the previous weight, we can calculate the weights recursively. So at round  $t$  we have a weight:

$$w_i^{(t)} = w_i^{(t-1)} \frac{\exp(-\alpha^{(t-1)} h^{(t-1)}(\mathbf{x}_i) y_i)}{2\sqrt{\varepsilon^{(t-1)}(1 - \varepsilon^{(t-1)})}} \quad (9)$$

Which is the same as:

$$w_i^{(t)} = w_i^{(t-2)} \frac{\exp(-\alpha^{(t-2)} h^{(t-2)}(\mathbf{x}_i) y_i)}{2\sqrt{\varepsilon^{(t-2)}(1 - \varepsilon^{(t-2)})}} \frac{\exp(-\alpha^{(t-1)} h^{(t-1)}(\mathbf{x}_i) y_i)}{2\sqrt{\varepsilon^{(t-1)}(1 - \varepsilon^{(t-1)})}} \quad (10)$$

Going all the way down to the base case, where  $w_1 = \frac{1}{n}$ , this becomes:

$$\begin{aligned} w_i^{(t)} &= \frac{1}{n} \frac{\prod_{\tau=1}^t \exp(-\alpha^{(\tau)} h^{(\tau)}(\mathbf{x}_i) y_i)}{\prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1 - \varepsilon^{(\tau)})}} \\ &= \frac{\exp(\sum_{\tau=1}^t -\alpha^{(\tau)} h^{(\tau)}(\mathbf{x}_i) y_i)}{n \prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1 - \varepsilon^{(\tau)})}} \end{aligned} \quad (11)$$

The decision function at round  $t$  is defined as:

$$f^{(t)} = \sum_{i=1}^t \alpha^{(i)} h^{(i)} \quad (12)$$

Substituting this into equation (11) we get:

$$w_i^{(t)} = \frac{\exp(-f^{(t)}(\mathbf{x}_i) y_i)}{n \prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1 - \varepsilon^{(\tau)})}} \quad (13)$$

### 1.3 Step 3

The goal is to prove for each  $i$  at any round  $t$

$$\frac{\exp(-f^{(t)}(\mathbf{x}_i)y_i)}{n \prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1-\varepsilon^{(\tau)})}} = \frac{\exp(-f^{(t)}(\mathbf{x}_i)y_i)}{\sum_{l=1}^n \exp(-f^{(t)}(\mathbf{x}_l)y_l)} . \quad (14)$$

The weights are normalized, that is:

$$\begin{aligned} \sum_{i=1}^n w_i^{(t)} &= 1 \\ \sum_{i=1}^n \frac{\exp(-f^{(t)}(\mathbf{x}_i)y_i)}{n \prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1-\varepsilon^{(\tau)})}} &= 1 \\ \sum_{i=1}^n \frac{1}{n} \sum_{i=1}^n \exp(-f^{(t)}(\mathbf{x}_i)y_i) \sum_{i=1}^n \frac{1}{\prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1-\varepsilon^{(\tau)})}} &= 1 \\ \sum_{i=1}^n \exp(-f^{(t)}(\mathbf{x}_i)y_i) &= \sum_{i=1}^n \prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1-\varepsilon^{(\tau)})} \\ \sum_{i=1}^n \exp(-f^{(t)}(\mathbf{x}_i)y_i) &= n \prod_{\tau=1}^t 2\sqrt{\varepsilon^{(\tau)}(1-\varepsilon^{(\tau)})} \quad (15) \end{aligned}$$

Substituting equation (15) into equation (13) we get:

$$w_i^{(t)} = \frac{\exp(-f^{(t)}(\mathbf{x}_i)y_i)}{\sum_{l=1}^n \exp(-f^{(t)}(\mathbf{x}_l)y_l)} . \quad (16)$$

## 2 Gradient Boosting

In gradient boosting, samples  $x_i$  are weighted at each round and have a larger weight if that sample was incorrectly predicted. This means different distributions of  $x_i$  are used in each round and this distribution depends on the error made on the previous round so the predicted  $\hat{y}_i^{(t)}$  are not independent from each other.

## 3 Recurrent Neural Networks

Replace the simple recurrent layer in `LSDA2020_RNN1.ipynb` notebook with a LSTM layer. Follow the equations from `LSDA2020_RNN2.ipynb` to implement the LSTM. You will need to add more variables and extend the step function. Also remember that the LSTM not only transfers the hidden state  $h_t$  to the next time step  $t+1$ , but also the cell state  $c_t$ .

The following parameters were added:

```
# LSTM parameters
# Forget layer parameters
W_f = tf.Variable(np.random.rand(dims, num_neurons), dtype=tf.float32)
b_f = tf.Variable(np.zeros((1, num_neurons)), dtype=tf.float32)
U_f = tf.Variable(np.random.rand(num_neurons, num_neurons), dtype=tf.float32)
# Input layer parameters
W_i = tf.Variable(np.random.rand(dims, num_neurons), dtype=tf.float32)
b_i = tf.Variable(np.zeros((1, num_neurons)), dtype=tf.float32)
U_i = tf.Variable(np.random.rand(num_neurons, num_neurons), dtype=tf.float32)
# Modulation layer parameters
W_j = tf.Variable(np.random.rand(dims, num_neurons), dtype=tf.float32)
```

```

b_j = tf.Variable(np.zeros((1, num_neurons)), dtype=tf.float32)
U_j = tf.Variable(np.random.rand(num_neurons, num_neurons), dtype=tf.float32)
# Output layer parameters
W_o = tf.Variable(np.random.rand(dims, num_neurons), dtype=tf.float32)
b_o = tf.Variable(np.zeros((1, num_neurons)), dtype=tf.float32)
U_o = tf.Variable(np.random.rand(num_neurons, num_neurons), dtype=tf.float32)
# Fully connected layer parameters
W_y = tf.Variable(np.random.rand(num_neurons, dims), dtype=tf.float32)
b_y = tf.Variable(np.zeros((1, dims)), dtype=tf.float32)

# Initial hidden state and cell state
h_0 = tf.Variable(np.zeros((1, num_neurons), dtype=np.float32))
c_0 = tf.Variable(np.zeros((1, num_neurons), dtype=np.float32))

# backpropagation
trainable_vars = [W_f, U_f, b_f, W_i, U_i, b_i, W_j, U_j, b_j, W_o, U_o, b_o]

```

The step function was changed to include all the LSTM gates:

```

@tf.function
def step(x_t, h, c):
    # forget layer
    f = tf.nn.sigmoid(
        tf.matmul(x_t, W_f) + tf.matmul(h, U_f) + b_f
    )
    # input layer:
    i = tf.nn.sigmoid(
        tf.matmul(x_t, W_i) + tf.matmul(h, U_i) + b_i
    )
    # modulation layer:
    j = tf.nn.tanh(
        tf.matmul(x_t, W_j) + tf.matmul(h, U_j) + b_j
    )
    # output layer:
    o = tf.nn.sigmoid(
        tf.matmul(x_t, W_o) + tf.matmul(h, U_o) + b_o
    )
    # memory (cell) layer:
    c = tf.math.multiply(f, c) + tf.math.multiply(i, j)
    # hidden state
    h = tf.math.multiply(o, tf.nn.tanh(c))
    # fully connected
    y_hat = tf.matmul(h, W_y) + b_y
    return y_hat, h, c

```

The `iterate_series` function was changed to include the cell state as an output and input.

```

@tf.function
def iterate_series(x, h, c):
    y_hat = []
    # iterate over time axis (1)
    for j in range(x.shape[1]):
        # give previous hidden state and input from the current time step
        y_hat_t, h, c = step(x[:, j], h, c)
        y_hat.append(y_hat_t)
    y_hat = tf.stack(y_hat, 1)
    return y_hat, h, c

```

The initial cell state and the modified `iterate series` functions were added to the training loop:

```

# train for a set number of iterations
for iteration in range(n_iterations):

```

```

# generates a long time series / normally loaded from dataset (e.g.
stocks, weather)
x, y = generateData(signal_length, predict_ahead, signal_repeats,
batch_size, noise_strength)
h = None
predictions = None
loss_list = []
grads = None

# do not feed complete series, but chunks of it (
truncated_backprop_length)
for i in range(0, total_series_length, truncated_backprop_length):

    with tf.GradientTape() as tape:
        if h is None:
            # initialize hidden state (h_0) -> new shape (batch_size,
num_neurons)
            h = tf.repeat(h_0, batch_size, 0)
            c = tf.repeat(c_0, batch_size, 0) #added
            x_part = x[:, i: i + truncated_backprop_length]
            y_part = y[:, i: i + truncated_backprop_length]

            # get predictions for this part (forward pass)
            y_hat, h, c = iterate_series(x_part, h, c) # changed

            # calculate mean squared error
            loss = tf.reduce_mean((y_hat - y_part)**2)
# backprop
            if grads is None:
                grads = tape.gradient(loss, trainable_vars)
            else:
                grads = grads + tape.gradient(loss, trainable_vars)

            loss_list.append(loss)
            # combine with previous predictions
            predictions = tf.concat([predictions, y_hat], 1) if predictions is
not None else y_hat

```

The mean training loss for the model was 0.003. Figure 1 shows the learned predictions.

## 4 Recurrent Neural Networks in Keras

Add different components to the RNN from `LSDA2020_RNN2.ipynb` and report the results on the validation set (the changed parts in the code).

1. Add bidirectional sequence processing by utilizing `tf.keras.layers.Bidirectional`.
2. Stack 2 LSTM layers. What is the difference to bidirectional processing? (you may need to use the `return_sequences` parameter)
3. Add a 1-d convolution layer (`tf.keras.layers.Conv1D`) before the recurrent part. You will need to reshape the data, you can use the `tf.keras.layers.Reshape` layer for that.
4. Gradient clipping can be helpful to train recurrent networks. Keras offers to clip gradients directly through the optimizer. Try this with clip values of 0.1, 1, and 10.

### 4.1 Bidirectional

```

def get_model_bidirectional(name="Bidirectional", shape=INPUT_SHAPE):
    inp = Input(shape)
    x = Bidirectional(LSTM(64))(inp)
    x = Dense(1)(x)

```




Figure 1: Learned Predictions of the LSTM model

```
model = Model(inp, x, name=name)
model.summary()
model.compile(sgd, loss='mae')
return model
```

Model: "Bidirectional"

Layer (type)	Output Shape	Param #
input_11 (InputLayer)	[(None, 19, 22)]	0
bidirectional_1 (Bidirection	(None, 128)	44544
dense_10 (Dense)	(None, 1)	129

Total params: 44,673

Trainable params: 44,673

Non-trainable params: 0

Model: Bidirectional

Mean Error: 1.9972226952917502

## 4.2 Stacked LSTM

```
def get_model_stacked(name="stacked", shape=INPUT_SHAPE):
    inp = Input(shape)
    x = LSTM(64, return_sequences=True)(inp)
    x = LSTM(64, go_backwards=True)(x)
    x = Dense(1)(x)
```

```

model = Model(inp, x, name=name)
model.summary()
model.compile(sgd, loss='mae')
return model

```

Model: "stacked"

Layer (type)	Output Shape	Param #
input_12 (InputLayer)	[(None, 19, 22)]	0
lstm_11 (LSTM)	(None, 19, 64)	22272
lstm_12 (LSTM)	(None, 64)	33024
dense_11 (Dense)	(None, 1)	65
Total params: 55,361		
Trainable params: 55,361		
Non-trainable params: 0		

Model: stacked

Mean Error: 1.9909247734882982

The bidirectional layers outputs two concatenated LSTM layers, while the stacked layers output one single LSTM layer. In the stacked layer, the input of the backward layer is affected by the forward layer since it passes through the forward layer, but the two layers are independent in the bidirectional layer.

### 4.3 Convolutional

```

def get_model_convolutional(name="convolutional", shape=INPUT_SHAPE):
    inp = Input(shape)
    x = Conv1D(filters=32, kernel_size=3, strides=1)(inp)
    x = LSTM(64, return_sequences=False)(x)
    x = Dense(1)(x)

    model = Model(inp, x, name=name)
    model.summary()
    model.compile(sgd, loss='mae')
    return model

```

Model: "convolutional"

Layer (type)	Output Shape	Param #
input_13 (InputLayer)	[(None, 19, 22)]	0
conv1d (Conv1D)	(None, 17, 32)	2144
lstm_13 (LSTM)	(None, 64)	24832
dense_12 (Dense)	(None, 1)	65

Total params: 27,041  
Trainable params: 27,041  
Non-trainable params: 0

-----  
Model: convolutional  
Mean Error: 2.0492551517846094

A recurrent neural network uses every piece of the input data to make a prediction. A convolutional neural network extracts features from the input data a subset of the input data and uses the features to make a prediction. Not all input data may be necessary to make a good prediction therefore convolutional networks can be less computationally expensive.

#### 4.4 Gradient Clipping

```
sgd = SGD(lr=0.01, momentum=0.9, nesterov=True, clipvalue = 1)
```

Model: RNN  
Mean Error: 2.0829751688270535